

COMP 251: Algorithms and Data Structures - Proofs Assignment

Liam Scalzulli

liam.scalzulli@mail.mcgill.ca

March 15, 2023

Complexity Proof

Claim: Inserting a node x into a red-black tree takes $O(\log n)$ time.

Presentation

Proof. To prove that inserting a node x into a red-black tree takes $O(\log n)$ time, we need to show that the number of operations performed during the insertion operation is proportional to the height of the tree, which is $O(\log n)$ as shown in the correctness proof.

The insertion operation in a red-black tree can be divided into three main steps:

1. **Binary search tree insertion:** This step involves finding the correct position for the new node x in the tree based on its key value, and inserting it as a regular binary search tree node.
2. **Coloring the new node red:** The new node is colored red to preserve the red-black tree properties.
3. **Fixing the red-black tree properties:** If the insertion of the new node causes a violation of the red-black tree properties, then one or more rotations and/or color flips are performed to restore the properties.

Let h be the height of the tree before the insertion operation, and let h' be the height of the tree after the operation. Since step 1 involves a standard binary search tree insertion, the number of comparisons performed during this step is proportional to the height of the tree, i.e., $O(h)$. Therefore, the worst-case time complexity of step 1 is $O(h)$, which is $O(\log n)$ in the average case for a balanced tree.

Step 2 involves a constant number of operations, namely setting the color of the new node to red. Therefore, the time complexity of this step is $O(1)$.

Step 3 involves fixing any violations of the red-black tree properties that may have occurred due to the insertion of the new node. The number of operations required to fix these violations is proportional to the height of the subtree affected by the violation. Since the subtree height is at most $h+1$, the worst-case time complexity of step 3 is $O(h+1)$, which is $O(\log n)$ in the average case for a balanced tree.

Therefore, the total worst-case time complexity of inserting a node into a red-black tree is $O(\log n)$. This is because step 1 takes $O(\log n)$, step 2 takes $O(1)$, and step 3 takes $O(\log n)$. Therefore, the total number of operations is proportional to the height of the tree, which is $O(\log n)$. \square

Source: Slides

Summary

The proof above shows that inserting a node x into a red-black tree takes $O(\log n)$ time. The insertion involves three steps: binary search tree insertion, coloring the new node red, and fixing any violations of the red-black tree properties. The time complexity of step 1 is $O(h)$, which is $O(\log n)$ in the average case for a balanced tree (as shown in the correctness proof). Step 2 takes $O(1)$ and step 3 takes $O(h + 1)$, which is also $O(\log n)$ in the average case. Therefore, the total worst-case time complexity of inserting a node into a red-black tree is $O(\log n)$.

Algorithm

Below you will find an algorithm that implements insertion for a red-black tree. The insertion algorithm includes the main three computational steps used in the proof, namely: BST insertion, coloring the new node red and fixing the red-black tree properties after the insertion was made. Many comments are made throughout the code to make things easier to understand. The full code includes a class *Node* which defines a red-black tree node, *Tree* which implements a small subset of red-black tree functionality and a *main* method for testing out the insertion.

```
/*
 * Definition for a red-black tree node.
 */
public class Node {
    int data;
    Node left, right, parent;
    boolean color;

    public Node(int data) {
        this.data = data;
    }

    public Node(int data, boolean color) {
        this.data = data;
        this.color = color;
    }
}

/*
 * A red-black tree with an insert operation.
 */
public class Tree {
    Node root;

    static final boolean RED = false, BLACK = true;

    /*
     * Default constructor.
     */
    public Tree() {}

    /*
     * Constructor with root.
     *
     * @param root The root of the tree.
     */
    public Tree(Node root) {
```

```

    this.root = root;
}

/*
 * Insert a node into the red-black tree.
 *
 * => Simulates a normal BST insertion but finishes
 * with a fixUp subroutine for violated RB-tree properties.
 *
 * @param key The data the new node should have.
 */
public void insert(int key) {
    Node node = root, par = null;

    // Depending on the key, we traverse left or right
    // => Same as in BST insertion.
    while (node != null) {
        par = node;
        if (key < node.data) node = node.left;
        else if (key > node.data) node = node.right;
        // Don't insert duplicate nodes
        else throw new IllegalArgumentException("Tree already contains a node with key " + key + ".")
    }

    // Color the new node red
    Node curr = new Node(key, RED);

    // Insert the node based on the key
    if (par == null) root = curr;
    else if (key < par.data) par.left = curr;
    else par.right = curr;

    // Set the parent of the newly inserted node
    curr.parent = par;

    // Fix the red-black tree properties
    fixUp(curr);
}

@Override
public String toString() {
    StringBuilder builder = new StringBuilder();
    toStringHelper(root, builder);
    return builder.toString();
}

/*
 * Fix red-black tree properties after an insertion.
 *
 * @param node The node to initiate the fix from.
 */
private void fixUp(Node node) {
    Node par = node.parent;

```

```

// Case 1: If the parent is null, we've reached the root, which is the end of the recursion.
// also, if the parent is black, there's nothing to do since the red-black properties are
// preserved.
if (par == null || par.color == BLACK) return;

// From this point on, the parent is red.
Node grandparent = par.parent;

//// Case 2: If the grandparent is null, that means the parent is the root.
// If we enforce black roots (rule 2), the grandparent will never be null.
// => Just need to recolor the parent node
//      R      B
//      \  ->  \
//      R      R
if (grandparent == null) {
    // As this method is only called on red nodes, all we have to do is recolor the root black.
    par.color = BLACK;
    return;
}

// Grab the uncle of the parent node
Node uncle = getUncle(par);

// Case 3: If the uncle is red, recolor the parent, grandparent, and uncle.
if (uncle != null && uncle.color == RED) {
    par.color = BLACK;
    grandparent.color = RED;
    uncle.color = BLACK;
    // Recursively fix the grandparent node
    fixUp(grandparent);
    // Parent is the left child of the grandparent.
} else if (par == grandparent.left) {
    // Case 4a: If the uncle is black and the node is the right child of its parent (forming a
    // left-right zigzag pattern).
    if (node == par.right) {
        rotateLeft(par);
        // Update the parent pointer to the new root node of the rotated subtree.
        par = node;
    }
    // Case 5a: If the uncle is black and the node is the left child of its parent (forming a
    // left-left straight pattern).
    rotateRight(grandparent);
    // Recolor the original parent (now the root of the subtree) and grandparent.
    par.color = BLACK;
    grandparent.color = RED;

    // Parent is the right child of the grandparent.
} else {
    // Case 4b: If the uncle is black and the node is the left child of its parent (forming a
    // right-left zigzag pattern).
    if (node == par.left) {
        rotateRight(par);
        // Update the parent pointer to the new root node of the rotated subtree.
        par = node;
    }
}

```

```

    }
    // Case 5b: If the uncle is black and the node is the right child of its parent (forming a
    // right-right straight pattern).
    rotateLeft(grandparent);
    // Recolor the original parent (now the root of the subtree) and grandparent.
    par.color = BLACK;
    grandparent.color = RED;
}
}

/*
 * Get the uncle node of a given input node.
 *
 * @param parent The node to grab the uncle from.
 * @return The uncle node.
 */
private Node getUncle(Node parent) {
    Node grandparent = parent.parent;
    if (grandparent.left == parent) return grandparent.right;
    if (grandparent.right == parent) return grandparent.left;
    throw new IllegalStateException("Parent is not a child of its grandparent.");
}

/*
 * Perform a right tree rotation.
 *
 * @param node The node to rotate from.
 */
private void rotateRight(Node node) {
    Node par = node.parent, left = node.left;
    node.left = left.right;
    if (left.right != null) left.right.parent = node;
    left.right = node;
    node.parent = left;
    replace(par, node, left);
}

/*
 * Perform a left tree rotation.
 *
 * @param node The node to rotate from.
 */
private void rotateLeft(Node node) {
    Node par = node.parent, right = node.right;
    node.right = right.left;
    if (right.left != null) right.left.parent = node;
    right.left = node;
    node.parent = right;
    replace(par, node, right);
}

/*
 * Replace the child of the parent node.
 *

```

```

    * @param par The parent node.
    * @param oldChild The old child of the parent.
    * @param newChild The new child to replace the old child with.
    */
private void replace(Node par, Node oldChild, Node newChild) {
    if (par == null) root = newChild;
    else if (par.left == oldChild) par.left = newChild;
    else if (par.right == oldChild) par.right = newChild;
    else throw new IllegalStateException("Node is not a child of its parent.");
    if (newChild != null) newChild.parent = par;
}

/*
 * A recursive helper method for building the toString
 * string builder.
 *
 * @param node The node to start from.
 * @param builder The string builder to append to.
 */
private void toStringHelper(Node node, StringBuilder builder) {
    builder.append(node.data);

    if (node.left != null) {
        builder.append(" L{");
        toStringHelper(node.left, builder);
        builder.append("}");
    }

    if (node.right != null) {
        builder.append(" R{");
        toStringHelper(node.right, builder);
        builder.append("}");
    }
}
}
}

```

Note: code adapted and repurposed from the following reference: <https://www.happycoders.eu/algorithms/red-black-tree-java/>

Code explanation

Below is the time graph for red-black tree insertion. Setting aside outliers and system differences (ran this on an M2 Pro Macbook 32GB RAM), it generally agrees with the assigned $O(\log n)$ time complexity.




Figure 1: Red-black tree insertion time graph

Real world example

Red-black trees are a type of binary search tree that automatically balance themselves. This makes them really useful in many computer science applications, including Database Management Systems (DBMS). When working with databases, it's important to have a good way to organize and search for information, which is where red-black trees come in.

One way red-black trees are used in DBMS is through something called a B+ tree. This is a special kind of tree that works really well with computer storage and is better for searching through lots of data. In a B+ tree, each node has a bunch of keys and pointers to other nodes, all sorted by the keys.

When adding new data to the tree, you need to find the right place for it. You start at the top of the tree and follow the pointers based on how the new data compares to the keys in each node. Once you find the right spot in a leaf node, you add the new data there. If the node is already full, you need to split it up and make a new node, then move the middle key up to the parent node so everything stays organized.

The cool thing about red-black trees (and B+ trees) is that they make sure the tree stays balanced, which means it takes less time to do things like search or add new data. Specifically, adding a new node to a red-black tree takes $O(\log n)$ time, where n is the number of entries in the tree. This is super important for big databases because it keeps everything running smoothly and quickly.

So, red-black trees are really helpful for organizing data in DBMS because they keep things balanced and make it easy to search, add, or delete data. They're the foundation for B+ trees, which are used a lot in DBMS to manage large databases efficiently.

References:

1. Cormen, T. H., Leiserson, C. E., Rivest, R. L., Stein, C. (2009). Introduction to Algorithms (3rd ed.). MIT Press.
2. Ramakrishnan, R., Gehrke, J. (2003). Database Management Systems (3rd ed.). McGraw-Hill.
3. Bayer, R., McCreight, E. M. (1972). Organization and Maintenance of Large Ordered Indices. Acta Informatica, 1(3), 173-189.

Correctness Proof

Claim: Red-Black Trees [CLRS 308]: A red-black tree with n internal nodes has height at most $2 \log(n + 1)$.

Presentation

Proof. We start by showing that the subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes. We prove this claim by induction on the height of x . If the height of x is 0, then x must be a leaf ($T.nil$), and the subtree rooted at x indeed contains at least $2^{bh(x)} - 1 = 2^0 - 1 = 0$ internal nodes. For the inductive step, consider a node x that has positive height and is an internal node with two children. Each child has a black-height of either $bh(x)$ or $bh(x) - 1$, depending on whether its color is red or black, respectively. Since the height of a child of x is less than the height of x itself, we can apply the inductive hypothesis to conclude that each child has at least $2^{bh(x)-1} - 1$ internal nodes. Thus, the subtree rooted at x contains at least $(2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$ internal nodes, which proves the claim. \square

Source: CLRS

Summary

The proof above uses induction to show that the subtree rooted at any node x in a red-black tree contains at least $2^{bh(x)} - 1$ internal nodes, where $bh(x)$ is the black-height of x . The base case is when x is a leaf, and the inductive step applies the inductive hypothesis to the two children of x and uses some arithmetic to derive the desired lower bound.

Algorithm

Below is an algorithm for computing the height of a red-black tree (same as computing the height for a BST), note that I use the same code as above and these methods are present in the provided *Tree* class.

```

/*
 * Compute the height of this red-black tree.
 *
 * @return The height of the tree.
 */
public int height() {
    return heightHelper(root);

/*
 * Compute the height of the given tree.
 *
 * @param The node to start from.
 * @return The height of the tree.
 */
private int heightHelper(Node node) {
    if (node == null) return 0;
    return 1 + Math.max(heightHelper(node.left), heightHelper(node.right));
}

```


Code explanation

Real world example

A real-world example where this height property is important is in search engines. Search engines have to manage tons of data and be able to quickly search through it to find what users are looking for. One way they do this is by using something called an inverted index.

An inverted index is a data structure that basically flips the relationship between documents and the terms they contain. Instead of listing the terms that appear in each document, an inverted index lists the documents in which each term appears. This makes it easier and faster to search for specific terms.

Red-black trees can be used to create an inverted index by storing the terms as keys and the lists of documents as values. When a search engine needs to find documents containing a specific term, it can search the red-black tree to find the term and get the list of documents quickly.

The height property of red-black trees (at most $2 \log(n + 1)$) is really important for search engines because it means that searching for a term in the tree takes a relatively small number of steps, even when there are lots of terms in the index. This helps search engines return results to users quickly, even when they're dealing with huge amounts of data.

So, red-black trees are great for search engines because they help create an efficient inverted index, which makes searching through massive amounts of data fast and easy. The height property of red-black trees ensures that the tree stays balanced and operations like searching, inserting, and deleting stay fast, which is super important when you're dealing with the crazy amounts of data search engines have to handle.

References:

1. Cormen, T. H., Leiserson, C. E., Rivest, R. L., Stein, C. (2009). Introduction to Algorithms (3rd ed.). MIT Press.
2. Manning, C. D., Raghavan, P., Schütze, H. (2008). Introduction to Information Retrieval. Cambridge University Press.
3. Knuth, D. E. (1998). The Art of Computer Programming, Volume 3: Sorting and Searching (2nd ed.). Addison-Wesley.